

NoSQL

In computing, **NoSQL** (sometimes expanded to "not only SQL") is a broad class of database management systems that differ from classic relational database management systems (RDBMSes) in some significant ways. These data stores may not require fixed table schemas, usually avoid join operations, and typically scale horizontally. Academia typically refers to these databases as **structured storage**,^{[1] [2] [3] [4]} a term that would include classic relational databases as a subset.

History

Carlo Strozzi used the term *NoSQL* in 1998 to name his lightweight, open-source relational database that did not expose an SQL interface.^[5] (Strozzi suggests that, as the current NoSQL movement "departs from the relational model altogether; it should therefore have been called more appropriately 'NoREL', or something to that effect."^[6])

Eric Evans, a Rackspace employee, reintroduced the term *NoSQL* in early 2009 when Johan Oskarsson of Last.fm wanted to organize an event to discuss open-source distributed databases.^[7] The name attempted to label the emergence of a growing number of non-relational, distributed data stores that often did not attempt to provide ACID (atomicity, consistency, isolation, durability) guarantees, which are the key attributes of classic relational database systems such as IBM DB2, MySQL, Microsoft SQL Server, PostgreSQL, Oracle RDBMS, Informix, Oracle Rdb, etc.

In 2011, work began on UnQL (Unstructured Query Language), a specification for a query language for NoSQL databases.^[8] It is built to query collections (versus tables) of documents (versus rows) with loosely defined fields (versus columns). So it is a superset of SQL where SQL is a very constrained type of UnQL where the queries will always return the same fields (same number, names and types). However, UnQL does not cover the DDL SQL statements like `CREATE TABLE` or `CREATE INDEX`^[9].

Architecture

Typical modern relational databases have shown poor performance on certain data-intensive applications, including indexing a large number of documents, serving pages on high-traffic websites, and delivering streaming media.^[10] Typical RDBMS implementations are tuned either for small but frequent read/write transactions or for large batch transactions with rare write accesses. NoSQL, on the other hand, can service heavy read/write workloads.^[10] Real-world NoSQL deployments include Digg's 3 TB for green badges (markers that indicate stories upvoted by others in a social network)^[11] and Facebook's 50 TB for inbox search.^[12]

NoSQL architectures often provide weak consistency guarantees, such as eventual consistency, or transactions restricted to single data items. Some systems, however, provide full ACID guarantees in some instances by adding a supplementary middleware layer (e.g., AppScale and CloudTPS).^{[13] [14]} Two systems have been developed that provide snapshot isolation for column stores: Google's Percolator system based on BigTable,^[15] and a transactional system for HBase developed at the University of Waterloo.^[16] These systems, developed independently, use similar concepts to achieve multi-row distributed ACID transactions with snapshot isolation guarantee for the underlying column store, without the extra overhead of data management, middleware system deployment, or maintenance introduced by the middleware layer.

Several NoSQL systems employ a distributed architecture, with the data held in a redundant manner on several servers, often using a distributed hash table. In this way, the system can readily scale out by adding more servers, and failure of a server can be tolerated.^[17]

Some NoSQL advocates promote very simple interfaces such as associative arrays or key-value pairs. Other systems, such as native XML databases, promote support of the XQuery standard. Newer systems such as CloudTPS also support join queries.^[18]

Taxonomy

NoSQL implementations can be categorized by their manner of implementation:

Document store

Name	Language	Notes
BaseX	Java, XQuery	XML database
Apache CouchDB	Erlang	
eXist	XQuery	XML database
Jackrabbit	Java	
Lotus Notes	LotusScript, Java, others	MultiValue
MarkLogic Server	XQuery	XML database
MongoDB	C++	BSON (Binary format JSON)
OrientDB	Java	
SimpleDB	Erlang	
Terrastore	Java	

Graph

Name	Language	Notes
AllegroGraph	SPARQL	RDF GraphStore
DEX	Java	High-performance Graph Database
InfiniteGraph	Java	High-performance, scalable, distributed Graph Database
Neo4j	Java	
OrientDB	Java	
FlockDB	Scala	
Sones GraphDB	C#	Graph database with query language called <i>GraphQL</i>
Pregel		

Key-value store

Key-value stores allow the application to store its data in a schema-less way. The data could be stored in a datatype of a programming language or an object. Because of this, there is no need for a fixed data model.^[19] The following types exist:

Eventually-consistent key-value store

- Apache Cassandra
- Dynamo
- Hibari
- Project Voldemort
- Riak^[20]

Hierarchical key-value store

- GT.M

Hosted services

- Freebase

Key-value cache in RAM

- Citrusleaf database
- memcached
- Oracle Coherence
- Redis
- Tuple space
- Velocity

Key-value stores on disk

- BigTable
- CDB
- Citrusleaf database
- Keyspace
- LevelDB
- membase
- Memcachedb
- Redis
- Tokyo Cabinet
- TreapDB
- Tuple space
- MongoDB

Ordered key-value stores

- Berkeley DB
- IBM Informix C-ISAM
- Memcachedb
- NDBM

Multivalue databases

- Extensible Storage Engine (ESE/NT)
 - OpenQM
 - Revelation Software's OpenInsight
 - Rocket U2
 - D3 Pick database
 - InterSystems Caché
-

Object database

- db4o
- GemStone/S
- InterSystems Caché
- JADE
- ObjectDB
- Objectivity/DB
- ObjectStore
- Versant Object Database
- ZODB

Tabular

- BigTable
- Apache Hadoop
- Apache Hbase
- Hypertable
- Mnesia

Tuple store

- Apache River

References

- [1] Hamilton, James (3 November 2009). "Perspectives: One Size Does Not Fit All" (<http://perspectives.mvdirona.com/CommentView,guid,afe46691-a293-4f9a-8900-5688a597726a.aspx>). . Retrieved 13 November 2009.
- [2] Lakshman, Avinash; Malik, Prashant. *Cassandra — A Decentralized Structured Storage System* (<http://www.cs.cornell.edu/projects/ladis2009/papers/lakshman-ladis2009.pdf>). Cornell University. . Retrieved 13 November 2009.
- [3] Chang, Fay; Jeffrey Dean, Sanjay Ghemawat, Wilson C. Hsieh, Deborah A. Wallach, Mike Burrows, Tushar Chandra, Andrew Fikes, and Robert E. Gruber. *Bigtable: A Distributed Storage System for Structured Data* (<http://labs.google.com/papers/bigtable-osdi06.pdf>). Google. . Retrieved 13 November 2009.
- [4] Kellerman, Jim. "HBase: structured storage of sparse data for Hadoop" (http://www.rapleaf.com/pdfs/hbase_part_2.pdf). . Retrieved 13 November 2009.
- [5] Lith, Adam; Jakob Mattson (2010). "Investigating storage solutions for large data: A comparison of well performing and scalable data storage solutions for real time extraction and batch insertion of data" (<http://publications.lib.chalmers.se/records/fulltext/123839.pdf>) (PDF). Göteborg: Department of Computer Science and Engineering, Chalmers University of Technology. p. 15. . Retrieved 2011-05-12. "Carlo Strozzi first used the term NoSQL in 1998 as a name for his open source relational database that did not offer a SQL interface[...]"
- [6] "NoSQL Relational Database Management System: Home Page" ([http://www.strozzi.it/cgi-bin/CSA/tw7/I/en_US/nosql/Home Page](http://www.strozzi.it/cgi-bin/CSA/tw7/I/en_US/nosql/Home%20Page)). Strozzi.it. 2007-10-02. . Retrieved 2010-03-29.
- [7] "NOSQL 2009" (http://blog.sym-link.com/2009/05/12/nosql_2009.html). Blog.sym-link.com. 2009-05-12. . Retrieved 2010-03-29.
- [8] <http://unqlspec.org/display/UnQL/Home>
- [9] Avram, Abel (04). "Interview: Richard Hipp on UnQL, a New Query Language for Document Databases" (<http://www.infoq.com/news/2011/08/UnQL>). <http://www.infoq.com>. . Retrieved 7 September 2011.
- [10] Agrawal, Rakesh et al. (2008). "The Claremont report on database research" (<http://db.cs.berkeley.edu/claremont/claremontreport08.pdf>). *SIGMOD Record* (ACM) **37** (3): 9–19. doi:<http://doi.acm.org/10.1145/1462571.1462573>. ISSN 0163-5808. .
- [11] "Looking to the future with Cassandra | Digg About" (<http://about.digg.com/blog/looking-future-cassandra>). About.digg.com. 2009-09-09. . Retrieved 2010-03-29.
- [12] "Cassandra" (http://www.facebook.com/note.php?note_id=24413138919&id=9445547199&index=9). facebook.com. 2008-08-25. . Retrieved 2011-08-19.
- [13] "Datastore Agnostic Transaction Support for Cloud Infrastructures" (<http://cs.ucsb.edu/~ckrintz/papers/ieecloud11.pdf>). IEEE. 2011-07-04. .
- [14] "CloudTPS: Scalable Transactions for Web Applications in the Cloud" (http://www.globule.org/publi/CSTWAC_ircs53.html). Globule.org. . Retrieved 2010-03-29.

- [15] "Large-scale Incremental Processing Using Distributed Transactions and Notifications" (http://www.google.ca/url?sa=t&source=web&cd=3&ved=0CCQQFjAC&url=http://www.usenix.org/events/osdi10/tech/full_papers/Peng.pdf&rct=j&q=Large-scale Incremental Processing Using Distributed Transactions and Notifications&ei=eM24TOYnjqedB_mHmLUN&usg=AFQjCNGGm1Xfaml5lq6Aj1R2BIX7WilIuQ&sig2=ZZcPWxhiMVSny-DmewlFIg&cad=rja). The 9th USENIX Symposium on Operating Systems Design and Implementation (OSDI 2010), Oct 4-6, 2010, Vancouver, BC, Canada. . Retrieved 2010-10-15.
- [16] "Supporting Multi-row Distributed Transactions with Global Snapshot Isolation Using Bare-bones [[HBase (<http://www.cs.uwaterloo.ca/~c15zhang/ZhangDeSterckGrid2010.pdf>)]]". The 11th ACM/IEEE International Conference on Grid Computing (Grid 2010), Oct 25-29, 2010, Brussels, Belgium. . Retrieved 2010-10-15.
- [17] "Cassandra: Structured Storage System over a P2P Network" (http://static.last.fm/johan/nosql-20090611/cassandra_nosql.pdf) (PDF). . Retrieved 2010-03-29.
- [18] "Consistent Join Queries in Cloud Data Stores" (http://www.globule.org/publi/CJQCDS_ircs68.html). Globule.org. . Retrieved 2011-01-31.
- [19] Marc Seeger (2009-09-21). "Key-Value Stores: a practical overview" (<http://dba.stackexchange.com/questions/607/what-is-a-key-value-store-database>). <http://www.slideshare.net/marc.seeger/keyvalue-stores-a-practical-overview>: slideshare. . Retrieved 2010-03-09. "Key value stores allow the application developer to store schema-less data. This data is usually consisting of a string that represents the key, and the actual data that is considered to be the value in the "key - value" relationship. The data itself is usually some kind of primitive of the programming language (a string, an integer, an array) or an object that is being marshalled by the programming languages bindings to the key value store. This replaces the need for fixed data model and makes the requirement for properly formatted."
- [20] "Riak: An Open Source Scalable Data Store" (<https://wiki.basho.com>). 28 November 2010. . Retrieved 28 November 2010.

External links

- (<http://www.odbms.org/downloads.aspx#nosql>) on [ODBMS.ORG: NoSQL Data Stores Section]
- NoSQLforums.ORG: NoSQL Knowledgebase - Live Message Board (<http://www.nosqlforums.org/>)
- NoSQL User Group (<http://www.linkedin.com/groups?gid=2085042>) on LinkedIn
- nosql-discussion (<http://groups.google.com/group/nosql-discussion>) on Google Groups
- nosqldatabases.com (<http://nosqldatabases.com/>)
- myNoSQL: news, articles and links about NoSQL (<http://nosql.mypopescu.com/>)
- nosql-databases.org (<http://nosql-databases.org/>)
- computerworld.com : No to SQL? Anti-database movement gains steam (http://www.computerworld.com/s/article/9135086/No_to_SQL_Anti_database_movement_gains_steam_)
- Is Microsoft Feeling the "NoSQL" Heat? (http://reddevnews.com/blogs/data-driver/2009/12/nosql-heat_0.aspx)
- *Information Week* "The NoSQL Alternative" (<http://www.informationweek.com/news/development/architecture-design/showArticle.jhtml?articleID=224900559>)
- How RDF Databases Differ from Other NoSQL Solutions (<http://blog.datagraph.org/2010/04/rdf-nosql-diff>)
- CouchOne (<http://www.couchone.com>)
- NoSql Tapes (<http://nosqltapes.com>)
- NoSQL Databases (Introduction and Overview) (<http://www.christof-strauch.de/nosql dbs.pdf>)

Article Sources and Contributors

NoSQL *Source:* <http://en.wikipedia.org/w/index.php?oldid=454406739> *Contributors:* Al3xpopescu, Alexandre.Morgaut, AlisonW, Amire80, Argv0, Arto B, Asafdapper, AxelBoldt, Bbulkow, Bdijkstra, Bearcat, Beland, Benatkin, Benhoyt, Bhaskar, Biofinderplus, Bovineone, CaptTofu, Ceefour, Cekli829, Charbelgereige, ChristianGruen, Clemwang, Cnorvell, ColdShine, Coldacid, Craigbeveridge, Crosbiesmith, Cybercobra, Cyril.wack, DamarisC, Dancrumb, DavidBourguignon, DavidSol, Davidhorman, Dericofilho, Dm, Dmccreary, Dmitri.grigoriev, Dredwolff, Drttm, Dshelby, Dstainer, Duncan, Ebalter, Eco schranzer, Edlich, Ehn, Eno, EricBloch, ErikHaugen, Ertugka, Euphoria, Excirial, Fiskbil, Fraktalek, Frap, Furrykef, Fxsjy, Gaborcsele, Germanviscuso, Getmoreatp, Gkorland, GlobalsDB, GoingBatty, Gpierre, Gstein, Heelmijnlevenlang, Hloeung, Hoelzro, Innortalnet, Irmatov, JLaTondre, Jabawack81, Jandalhandler, Javalangstring, Jeffdexter77, JnRouvignac, Jonasfagundes, Joolean, Jottinger, Jrudisin, Jstplace, Justinsheehy, Kgffleischmann, Ki2010, KiloByte, Kkbhumana, Koavf, Komap, Korrawit, Leotohill, Lfstevens, Lguzenda, Linas, Looris, Luisramos22, MMSequeira, Mabdul, Magnuschr, Marasmusine, Mbonaci, Mhegi, Miami33139, Mitpradeep, Mjresin, Morphh, Mortense, MrWerewolf, Mshefer, Mtrenceni, Mydoghasworms, Natishalom, Nawroth, Netmesh, Nileshbansal, Ntoll, Omidnoorani, PatrickFisher, Pcap, Peak, Phillips-Martin, Phunehehe, Plustgarten, Pnm, Poohneat, R39132, Rabihnassar, Really Enthusiastic, Rfl, RobertG, Robhughadams, Ronz, Rtweed1955, Russss, Sae1962, SamJohnston, ScottConroy, Sduplooy, Seancribbs, Seraphimblade, Shadowjams, Shepard, Shijucv, Smyth, Sorenriise, Sstrader, Stephen Bain, Stephen E Browne, Stevedekorte, Stimpy77, Syaskin, TJRC, Tagishsimon, Tedder, Theandrewdavis, Thomas.uhl, ThomasMueller, Thumperward, Tobiasivarsson, Tomdo08, Trbdavies, Tshanky, Tuvrotya, Uhbif19, Violaaa, Viper007Bond, Volt42, Voodootikigod, Vychtrle, Weimann, Whooym, William greenly, Winterst, Woohookitty, Wyverald, YPavan, Zapher67, Zond, 419 anonymous edits

License

Creative Commons Attribution-Share Alike 3.0 Unported
[//creativecommons.org/licenses/by-sa/3.0/](http://creativecommons.org/licenses/by-sa/3.0/)